

# Deploying Infrastructure for Dask on Google Cloud

**Sang-Yun Oh**

Dept. of Statistics and Applied Probability  
University of California Santa Barbara

# What is Dask?

*Dask breaks up large computations and route parts of them efficiently onto distributed hardware. Dask is routinely run on thousand-machine clusters ...*

## About Dask

- [Description of Dask](#)
- [Example distributed applications](#)
- [Distributed infrastructure](#)

# Computing Infrastructure for Dask

- Dask run on laptops, clusters, and HPC environments
- Develop/test in small scale (laptops)
- Deploy/scale in large scale (clusters and HPC environments)

## Which Infrastructure to Use?

- **Cloud computing**: slow(er) but easy access and on-demand
- **HPC**: highly optimized but access is restricted and may not be interactive

# Cloud Computing and Reproducibility

What is reproducibility?

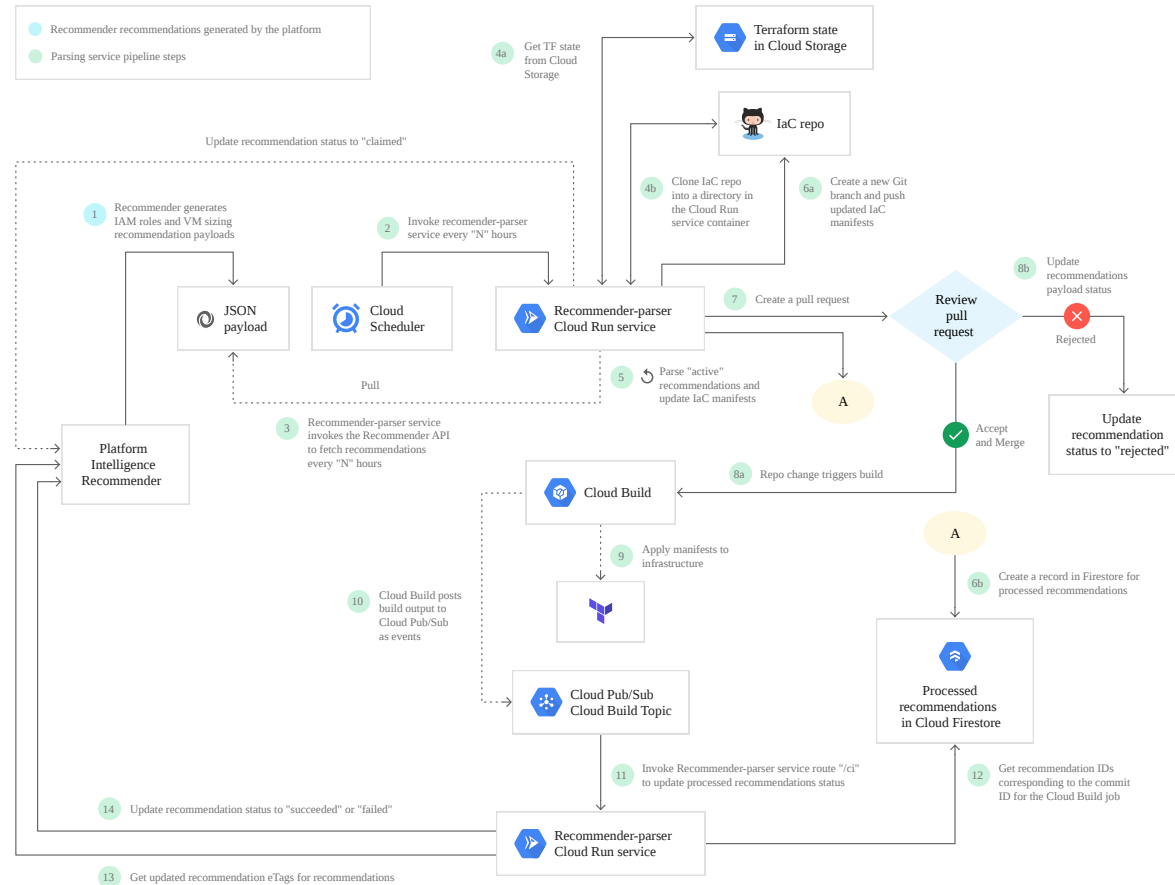
*An article about computational result is advertising, not scholarship. The actual scholarship is the **full software environment, code and data**, that produced the result.*

# Cloud Platforms and Infrastructure as Code (IaC)

- Cloud platforms have APIs
- 😊 Setups of cloud infrastructure can be coded (IaC)
- 👍 Analysis code and cloud IaC on GitHub + Data → **Reproducibility**

Unfortunately, cloud infrastructure is mostly Do-It-Yourself (DIY) 😞

# Infrastructure as Code

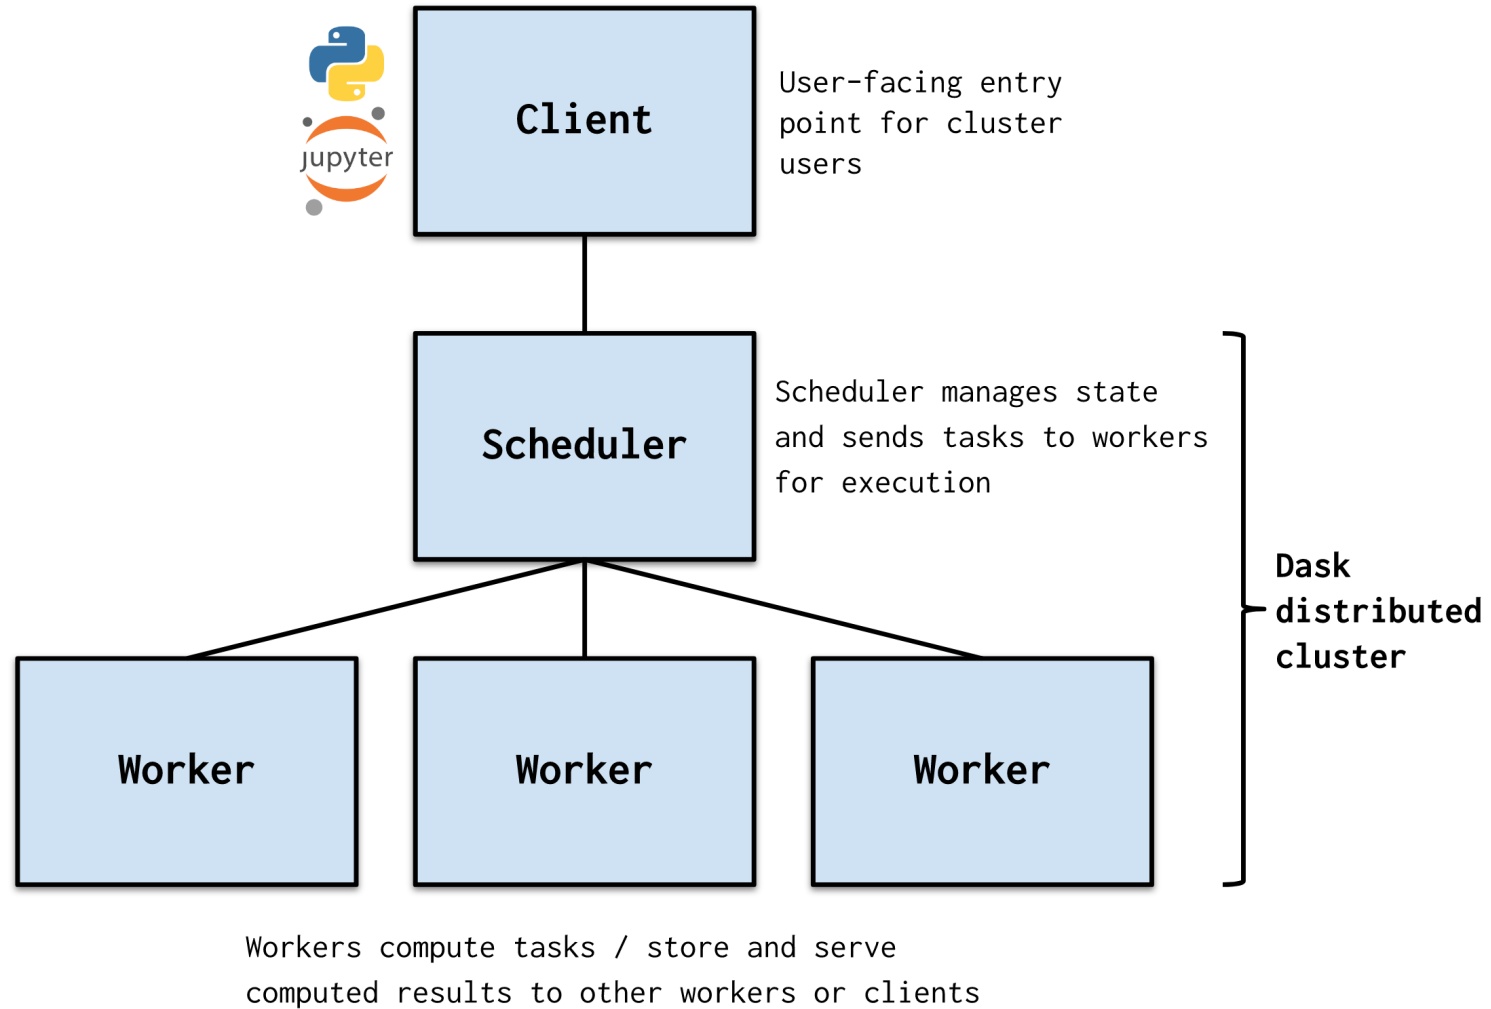


(Using Recommendations for Infrastructure as Code)

# Goal for Today

- Conceptual understanding of distributed computing infrastructure for Dask
- Discuss enabling technologies
- Hands-on lab using Google Cloud Platform

# Dask Architecture





# Architecture Components

Client, scheduler, and worker processes can be distributed in different ways

- 🏠 *Before 2006*: physical server run (a combination of) processes
- 🏡 *Before 2015*: multiple virtual machines on powerful machines
- 🏢 *After 2015*: Kubernetes cluster to handle orchestration


**What is Kubernetes?** 🧑

# Kubernetes and Containers

**Kubernetes** is an open-source system for automating deployment, scaling, and management of containerized applications.

**Application containerization** (e.g. **Docker**) is an OS-level virtualization method used to deploy and run distributed applications without launching an entire virtual machine (VM) for each app.

# Enabling Technologies



**docker**

Component  
contents



Components wiring  
diagram



**kubernetes**

Builder and  
manager



Google Cloud

Provides resources

# Hands-on Lab

Google Cloud Command Line Tool: <https://shell.cloud.google.com/>

## Main tools

```
gcloud --version      # controls Google Cloud resources
docker --version      # container-level controls
kubectl version       # cluster-level controls
helm version          # installation "blueprint"
```

# Login to Google Cloud

```
gcloud auth list          # check current account
# gcloud auth login syoh@ucsb.edu  # login using another account

# set default project
gcloud config set project testing-sandbox-324502
gcloud config set compute/zone us-central1-a
```

# Start Kubernetes Cluster

```
# create Kubernetes cluster
gcloud container clusters create \
  --machine-type e2-standard-4 \
  --num-nodes 2 \
  [unique-cluster-name]
```

- [Machine type documentation](#)
- [Regions and zones documentation](#)
- Use unique cluster names: e.g. include your initials

# Dask Cluster Helm Chart

Three important concepts for Helm:

1. **Chart**: a bundle of information necessary to create an instance of a Kubernetes application. (the blueprint)
2. **Config**: configuration information that can be merged into a packaged chart. (user specified setting)
3. **Release**: running instance of a *chart*, combined with *config*. (deployed instance)

## Dask Helm Chart

Blueprint for setting up *Jupyter Lab*, *scheduler*, and *workers* on Kubernetes cluster

# Dask Helm Chart Config

- User configuration supercedes default values in `values.yaml` file

```
cat << EOF > config.yaml
jupyter:
  serviceType: "LoadBalancer" # makes Jupyter notebook publicly accessible
scheduler:
  serviceType: "LoadBalancer" # makes Dask scheduler publicly accessible
worker:
  replicas: 4
EOF
```

- Instantiate Dask Cluster

```
helm repo add dask https://helm.dask.org/
helm repo update
helm install --wait my-dask -f config.yaml dask/dask # takes a while
```



# Kubernetes running **my-dask** release

```
sangoh@cloudshell:~ (testing-sandbox-324502)$ kubectl get all
```

NAME	READY	STATUS	RESTARTS	AGE
pod/my-dask-jupyter-54ddbffd9d-rbsjb	1/1	Running	0	8m45s
pod/my-dask-scheduler-7f4f94bb7d-4c4fn	1/1	Running	0	8m45s
pod/my-dask-worker-6877d8f79f-42jrb	1/1	Running	0	8m44s
pod/my-dask-worker-6877d8f79f-88wvm	1/1	Running	1	8m45s
pod/my-dask-worker-6877d8f79f-9qvc5	1/1	Running	0	8m44s
pod/my-dask-worker-6877d8f79f-g659d	1/1	Running	1	8m45s
pod/my-dask-worker-6877d8f79f-htg8p	1/1	Running	1	8m45s
pod/my-dask-worker-6877d8f79f-t44jn	1/1	Running	1	8m45s

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/kubernetes	ClusterIP	10.12.0.1	<none>	443/TCP	11m
service/my-dask-jupyter	LoadBalancer	10.12.8.243	34.121.167.99	80:32536/TCP	8m45s
service/my-dask-scheduler	LoadBalancer	10.12.14.160	34.67.163.35	8786:30322/TCP, 80:32005/TCP	8m45s

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/my-dask-jupyter	1/1	1	1	8m45s
deployment.apps/my-dask-scheduler	1/1	1	1	8m45s
deployment.apps/my-dask-worker	6/6	6	6	8m45s

NAME	DESIRED	CURRENT	READY	AGE
replicaset.apps/my-dask-jupyter-54ddbffd9d	1	1	1	8m46s
replicaset.apps/my-dask-scheduler-7f4f94bb7d	1	1	1	8m46s
replicaset.apps/my-dask-worker-6877d8f79f	6	6	6	8m46s

Google Cloud Kubernetes Workloads Overview [🔗](#)

# Server addresses

```
export DASK_SCHEDULER=$(kubectl get svc --namespace default my-dask-scheduler -o jsonpath='{.status.loadBalancer.ingress[0].ip}')
export DASK_SCHEDULER_UI_IP=$(kubectl get svc --namespace default my-dask-scheduler -o jsonpath='{.status.loadBalancer.ingress[0].ip}')
export DASK_SCHEDULER_PORT=8786
export DASK_SCHEDULER_UI_PORT=80

export JUPYTER_NOTEBOOK_IP=$(kubectl get svc --namespace default my-dask-jupyter -o jsonpath='{.status.loadBalancer.ingress[0].ip}')
export JUPYTER_NOTEBOOK_PORT=80

echo tcp://$DASK_SCHEDULER:$DASK_SCHEDULER_PORT          -- Dask Client connection
echo http://$DASK_SCHEDULER_UI_IP:$DASK_SCHEDULER_UI_PORT -- Dask dashboard
echo http://$JUPYTER_NOTEBOOK_IP:$JUPYTER_NOTEBOOK_PORT  -- Jupyter notebook
```

# Run Example Notebook

1. Open Jupyter notebook specified in output
2. Copy dashboard URL into Dask Jupyter lab extension
3. Open and run `examples/04-dask-array.ipynb`
4. Additional packages are needed for `examples/05-nyc-taxi.ipynb`



# Install Additional Packages

```
cat << EOF > config.yaml
jupyter:
  serviceType: "LoadBalancer" # makes Jupyter notebook publicly accessible
  env:
    - name: EXTRA_PIP_PACKAGES
      value: "pyarrow gcsfs"

scheduler:
  serviceType: "LoadBalancer" # makes Dask scheduler publicly accessible

worker:
  replicas: 4
  env:
    - name: EXTRA_PIP_PACKAGES
      value: "pyarrow gcsfs"

EOF
```

# Upgrade `my-dask` Release

- `upgrade` release rather than `install`

```
helm upgrade --wait my-dask -f config.yaml dask/dask # takes a while
```

- Run `examples/05-nyc-taxi.ipynb` (will break). Check why with `kubectl`

```
kubectl get all # what do you notice?  
kubectl describe [pod/my-dask-worker-000000] # evicted resource name
```

# Try again

```
cat << EOF > config.yaml
jupyter:
  serviceType: "LoadBalancer" # makes Jupyter notebook publicly accessible
  resources:
    limits:
      cpu: 1
      memory: 3G
    requests:
      cpu: 0.5
      memory: 2G
  env:
    - name: EXTRA_PIP_PACKAGES
      value: "pyarrow gcsfs matplotlib"

scheduler:
  serviceType: "LoadBalancer" # makes Dask scheduler publicly accessible
  resources:
    limits:
      cpu: 1
      memory: 3G
    requests:
      cpu: 0.5
      memory: 2G

worker:
  replicas: 15
  resources:
    limits:
      cpu: 1
      memory: 3G
    requests:
      cpu: 0.5
      memory: 2.5G
  env:
    - name: EXTRA_PIP_PACKAGES
      value: "pyarrow gcsfs"

EOF
```

# Upgrade `my-dask` Release again

- Resize cluster to add nodes

```
gcloud container clusters resize [your-cluster-name] --num-nodes 5
```

- `upgrade` release rather than `install`

```
helm upgrade --wait my-dask -f config.yaml dask/dask # takes a while
```

- Run `examples/05-nyc-taxi.ipynb`